


A Circus of Circuitry – A COMPI521 Notation

THEORY

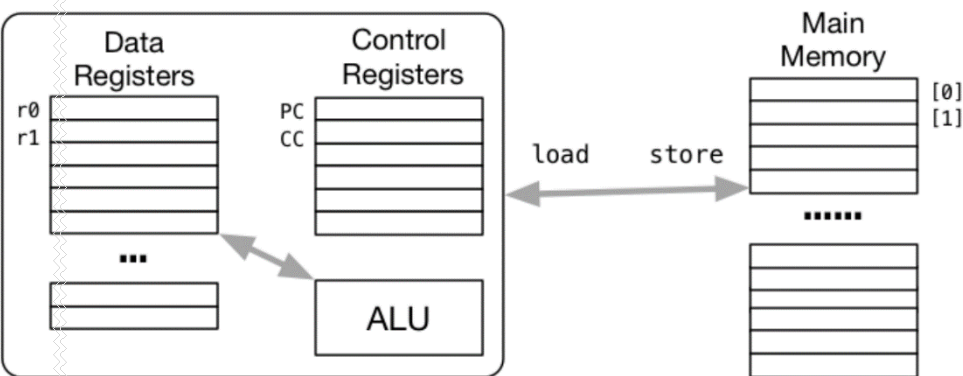
CPUs – Level 9

- **Processors (CPU)**
 - Control, arithmetic, logic, bit operators
 - Small set of simple operations
 - A range of ways to determine data locations
 - Small amount of very fast storage (registers)
 - Small number of control registers
 - Fast fetch-decode-execute cycle
 - Access to system bus to communicate with other components
- **Storage**
 - Memory (small and fast)
 - Very large random-addressable array of bytes
 - Can fetch single bytes, or multiple-byte chunks into CPU
 - Access time typically $0.1\mu s$, size perhaps 64GB
 - Disk storage (large and slower)
 - Very very large block-oriented storage
 - Often on a spinning disk, 512-4KB per request
 - Access time typically 30ms, size 8TB
 - Access time nowadays $100\mu s$, size 512GB
- **CPU Architecture** – modern processors (CPUs) have:
 - A set of data registers
 - Fast storage used to contain data while it is being manipulated
 - A set of control registers
 - Places to control the actions of a CPU
 - Including PC – Program Counter
 - A control register which keeps track of execution
 - An arithmetic-logic unit (ALU)
 - Access to memory (RAM)
 - A set of simple instructions to:
 - Transfer data between memory and registers
 - Push values through the ALU to compute results



- Make tests and transfer control of execution

CPU



- Execution logic looks like the right

```
while (1)
{
    instruction = memory[PC]
    PC++ // move to next instr
    if (instruction == HALT)
        break
    else
        execute(instruction)
}
```

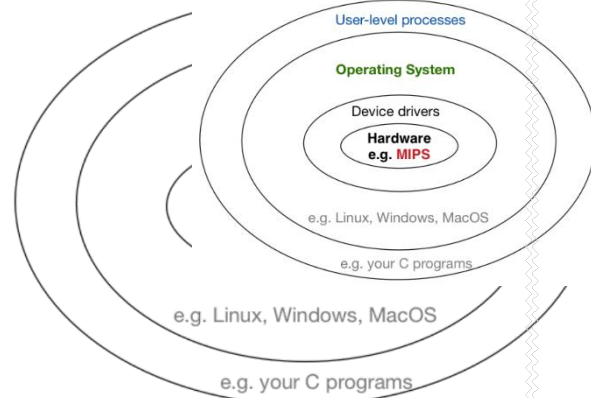
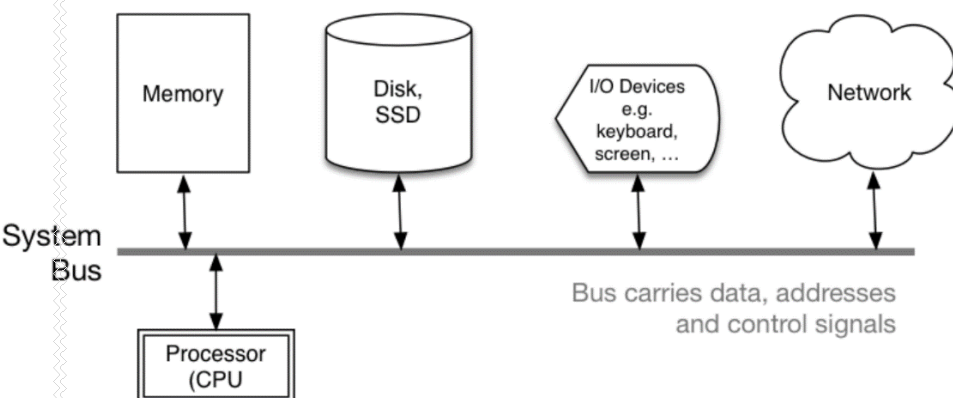
Instructions

- They are stored as strings of bits (as all things ultimately are)
 - 32-bit bit-strings
- They contain info about:
 - What the **operator** is
 - What **registers** are involved
 - What **memory locations** are involved
- They carry out the operation
- They store the value (if required) in memory



Address	Content
0x100000	0x3c041001
0x100004	0x34020004
0x100008	0x0000000c
0x10000C	0x03e00008

Computer Architecture – *Like an Abacus but Amazing*




- Computer system**
 - Hardware core surrounded by layers of software
- Modern computer**
 - Devices connected via a system “bus”

- Bus – a communication system that transfers data between the components of a computer

- **Operating systems (OSs)**

- Provide an abstraction layer on top of hardware (provide the same view regardless of underlying hardware)
- Have privileged access to the raw machine
- Manage the use of machine resources (CPU, disk, memory...)
- Provide uniform interface to access machine-level operations
- Arrange for controlled execution of user programs
- Provide multi-tasking and parallelism

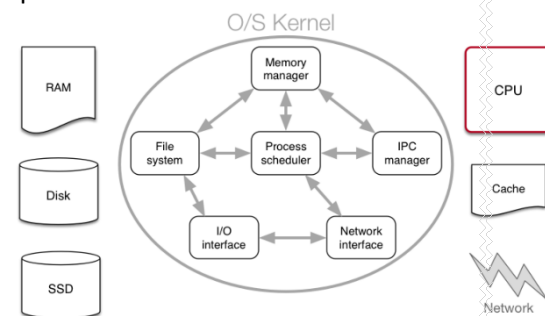


- **Operating systems 'flavours'**

- Batch (eg Eniac, early IBM OSs)
 - Computational jobs run one-at-a-time in a queue
- Multi-user (eg Unix/Linux, OSX, Windows)
 - Multiple jobs appear to run parallel
- Embedded (eg Android, iOS)
- Small, cut-down OS embedded in a device
- Real-time (eg RTLinux)
 - Specialised OS with time guarantees on job completion

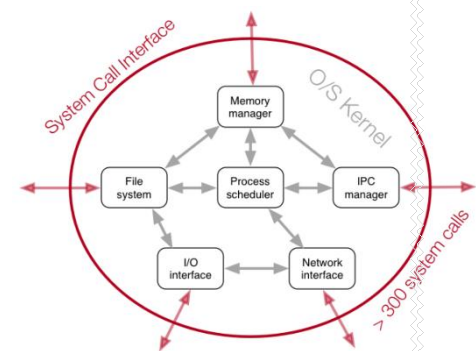
- **Modern OSs abstractions**

- Users – who can access the system
- Access rights – what users are allowed to do
- File system – how data is organised
- Input/output – transferring data to/from devices
- Processes – active computation on the system
- Communication – how processes interact
- Networking – how the system talks to other systems
- Core OS functions stem from the operating system **kernel**



- **Execution modes**

- Privileged mode
 - Full access to all machine operations and memory regions
- Non-privileged (user) mode
 - A limited set of operations (still Turing complete)
 - Access to only part of the memory
- System calls allow programs to cross the privileged/user boundary in a controlled manner via well-defined requests
 - Full OSs provide 100s of system calls
 - Process management
 - File management
 - Device management
 - Information maintenance
 - Communication



- **Libraries**


- We resort to these when we want to do something beyond the low-level operations of system calls
- Collections of useful functions

- Referenced from within user programs (eg C functions)
- Defined by #include <xxx.h>
- Integrated with user code at 'link time'
- **Applications**
 - User-level programs which perform some useful task
 - Possibly supplied with system (ls, vim, gcc)
 - Possibly implemented by users (dcc, check, Webcms3)
 - They live in /bin, etc
 - Generally built using libraries, may also use system calls
 - Unix is unusual
 - Has a command interpreter that runs as a user-level process
 - But can invoke other user-level processes
- **System calls**
 - Invoked directly (library provided in the manual for the OS)
 - Invoked indirectly through functions in the C libraries (also documented in the manual)
 - May fail
 - This can be detected by checking return values or the global variable **errno**
 - C programs need to check and handle error themselves (provides no exception handling)
- **Failed calls in C**
 - fprintf(stderr, "Can't do %s", *Something*);
exit(1)
 - perror(char **Message*);
exit(errno); (sends precise error to shell)
 - error(*Status*, *ErrNum*, *Format*, *Expressions*, ...); (Linux library function)
 - Print error message using program name, *Format* and *Expressions*
 - If *Status* is non-zero, invoke exit(*Status*) after printing message
 - If *ErrNum* is non-zero, also print standard system error message

File Systems – *The Bureaucracy*

- **File systems**
 - Provide a mechanism for managing stored data
 - Typically on a disk device
 - Allocates chunks of space on the device to files
 - A file can be viewed as a sequence of bytes
 - A directory can be viewed as a file containing references to other files
 - Allows access to files by name and with access rights
 - Arranges access to files via directories (folders)
 - Maintains information about files/directories (meta-data)
 - Deals with damage on the storage device ("bad blocks")

- Unix/Linux file system
 - “Tree structured”
 - Actually a graph after considering symbolic links
 - Used to access various types of objects
 - Files, directories, devices, processes, sockets
 - Paths can be
 - Absolute (full path from root)
 - Relative (path starts from CWD)
- File manipulation commands
 - FILE * and it’s relevant operations
 - fgets, fputs, fgetc, fputc
 - fscanf, fprintf
- **Unix file system** defines a range of file-system-related types
 - **off_t**
 - Offsets within files
 - Typically long and signed to allow backward references
 - **size_t**
 - Number of bytes in some object
 - (Basically an unsigned int or long long)
 - Unsigned
 - **ssize_t**
 - Sizes of read/written blocks
 - Like size_t, but signed to allow for error values
 - **struct stat**
 - File system object metadata
 - Stores info *about* a file
 - Requires ino_t, dev_t, time_t, uid_t
- **Inodes**
 - Data structures which store metadata for file system objects
 - Physical location on storage device
 - File type, file size
 - Ownership, access permissions, timestamps
 - Each file system volume (designated storage area) has a table of inodes
 - Access to a file by name requires a directory
 - Directories are lists of name/inode pairs
 - Access to files looks like this:
 - Open directory and scan for name
 - If not found, “No such file or directory”
 - If found (name, ino) access inode table inodes[ino]
 - Collect file metadata
 - Check file permissions
 - If no permission, “Permission denied”
 - Collect information, update timestamp
 - Use physical location to access device and manipulate file data
- Unix presents a uniform interface to file system objects
 - Functions and syscalls manipulate objects as a stream of bytes
 - Accessed via a file descriptor (index into a system table)



- Common operations are open(), close(), read(), write(), etc
- **int open(char *Path, int Flags)**
 - Attempt to open an object at Path according to Flags
 - Flags are constants defined in <fcntl.h>
 - O_RDONLY (open for reading)
 - O_WRONLY (open for writing)
 - Etc
 - They can be combined
- **int close(int FileDesc)**
 - Attempt to release an open file descriptor
 - If this is the last reference to object, release its resources
- **ssize_t read(int FileDesc, void *Buffer, size_t Count)**
 - Attempt to read Count bytes from FileDesc into Buffer
 - Does not check whether Buffer has enough space
 - Returns number of bytes read
- **ssize_t write(int FileDesc, void *Buffer, size_t Count)**
 - Attempt to write Count bytes from Buffer onto FileDesc
 - Does not check whether Buffer has Count bytes of data
 - Returns number of bytes written
 - Useful – can write into structs
- **off_t lseek(int FileDesc, off_t Offset, int Whence)**
 - Set the “current position” of the FileDesc (basically we move around in a file)
 - Offset is in units of bytes, can be negative
 - Whence can be:
 - SEEK_SET – set to Offset from start of file
 - SEEK_CUR – set to Offset from current position
 - SEEK_END – set file position to Offset from end of file
- **int stat(char *FileName, struct stat *StatBuf)**
 - Stores metadata associated with FileName into StatBuf
 - Includes
 - Inode number, file type and access mode, owner, group, etc
 - Size in bytes, storage block size, allocated blocks
 - Time of last access, modification, status change
 - **int fstat(int FileDesc, struct stat *StatBuf)**
 - Same but gets data via an open file descriptor
 - **int lstat(char *FileName, struct stat *StatBuf)**
 - Same but doesn’t follow symbolic links
- Links
 - Hard links
 - Multiple directory entries referencing the same inode
 - The two entries must be on the same filesystem
 - Symbolic links (symlinks)
 - A file containing the path name of another file
 - Opening the symlink opens the file being referenced
 - **int mkdir(char *PathName, mode_t Mode)**
 - Create a new directory called PathName with mode Mode


- All files in the path must exist
- The directory we write to must be writeable to the caller
- The directory we create must not already exist
- The new directory contains two initial entries
 - . is a reference to itself
 - .. is a reference to its parent directory
- Returns 0 if successful, otherwise -1
- **int fsync(int FileDesc)**
 - Ensures that data associated with FileDesc is written to storage
 - (Usually written data is stored in memory buffers, and eventually goes to permanent storage)
 - (fsync()) forces this to happen immediately

Processors and Processing – *It's a Process*

- A process is an active computation, consisting of:
 - RAM: code (read-only), data (read/write)
 - Registers: program counter (PC) and other registers
 - Other management info
- Multiple processes can be active simultaneously
 - Typically not all loaded in RAM at once
 - Processes can be suspended (waiting)
 - Restoring a process: load code, data, registers

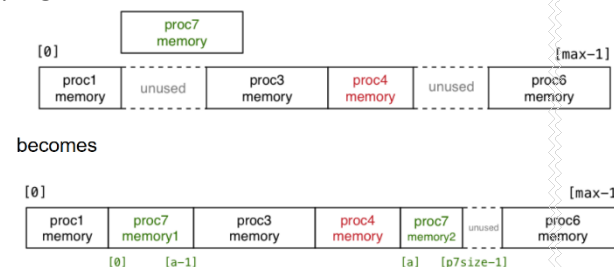
- If we have multiple processes...

- We can't structure our program linearly, since the stack frames would interfere with each other
- We thus allocate some memory for both processes
- When we reference memory in our code, we can't do it directly, so we need to remember the base address for each process
- We then reference specific areas of memory with $addr + p1size$ instead of $addr$
 - This can be automatically mapped during execution
- Remember base address for each process (process table)
 - When process starts, load base into mapping hardware
 - Interpret every $addr$ as $base + addr$



- If we add a new process...

- Slot it in a free chunk of memory that's big enough
- Otherwise, split our process into two smaller chunks of memory




- **Memory Management**

- It makes sense to standardise chunks of memory
 - Call each chunk of address space a **page**
 - All pages are the same size **P**
 - Page **i** has addresses **A** in the range $i * P \leq A < (i + 1) * P$
- **Virtual addresses** and address mapping tables
 - Each process has an array of page entries
 - Each page entry contains the start address of one chunk
 - Can compute the index of relevant page entry by A/P
 - Can compute the offset within the page by $A \% P$
- Eg: Pages p0@5000, p1@3000, p2@1000, page size 1000
 - Virtual address 0 refers to physical address 5000
 - Virtual address 128 refers to physical address 5128
 - Virtual address 1500 refers to physical address 3500
 - Virtual address 2200 refers to physical address 1200
- If pages are of size $Pagesize == 2^n \dots$
 - Computing page number and offset becomes very efficient
 - We simply use bitwise operators with two bitmasks (for pageno and offset)

Virtual Memory

- A side effect of this mapping is that we don't need to load all of our process's pages upfront (like we would in MIPS)
 - We start with a small memory footprint (ie just main and stack top)
 - We load new process address pages into memory as needed
 - Grow up to the size of the physical memory
- We can:
 - Divide process memory space into fixed-size pages
 - Load process pages into physical memory
- This is virtual memory
 - We map virtual to physical memory in a table




- Implications
 - In any given time, a process is likely to access only a fraction of its pages
 - This is known as the **working set model**
 - We only need to hold the working set at a given time
 - If each process has a relatively small working set, it can hold pages for many processes simultaneously
 - *The process address space can be larger than immediate physical memory*

Where do we "load" pages from?

- Code, global data from executable file stored on disk
- Dynamic (heap, stack) data is created in memory
- Note: transferring between disk → memory is very expensive

Pre-process page table

- Can be considered a struct of structs, stored in memory
- Allows us to keep track of what is and is not loaded
- Requesting a non-loaded page generates a **page fault**
 - We often keep a **free list** to keep track of free frames
 - If there are no free frames, we either *suspend* the request or *replace* a currently loaded page



Page K-2
not yet
created
on disk


- Page replacement
 - When a page is replaced
 - If it's been modified since loading, save to disk (requires flags)
 - Grab its frame number and give it to the requestor
 - Which frame should be replaced
 - Define a “usefulness” measure for each frame
 - Grab the frame with the lowest usefulness. When replacing...
 - Prefer pages that are read-only (no need to write to disk)
 - Prefer pages that are unmodified (no need to write to disk)
 - Prefer pages that are only used by one process
 - “A page not used recently may not be needed again soon” (LRU replacement)
 - **Thrashing**: constantly swapping pages in and out of memory
 - The working set model and LRU helps avoid this

- Page replacement strategies

- **LRU**
 - Least recently used
- **FIFO**
 - Page frames are entered into a queue and page replacement uses the frame from the front of the queue
- **Clock sweep**
 - Uses a reference bit for each frame
 - Maintains a circular list of allocated frames
 - Uses a “clock hand” which iterates over page frame list
 - All referenced pages are skipped and reset
 - First-found unreferenced frame used

- Page table strategies (we only need in practice some $n \ll K$ page table entries at one time)

- Hashing
 - Make a page table
 - Hash it (map several pages to one spot to save space)
- Multi-level page tables
 - One table of tables
 - I.e entries 0 to 7 in one table, 8 to 15 in another, etc
 - It's unlikely we use all at once



- **Cache memory**

- Small, fast memory close to CPU
- Holds parts of RAM we think will be heavily used
- Transfers data to/from RAM in blocks
- Memory reference hardware checks cache first
 - If not there, gets from RAM
- Similar replacement strategies similar to virtual memory

- Memory management hardware

- Address translation is done by MMU
- TLB – translation lookaside buffer

- **Process**

- An instance of an executing program

- Each process has an execution state defined by
 - Current execution point
 - Current values of CPU registers
 - Current contents of virtual address space
 - Information about open files, sockets, etc
- To manage processes, maintain
 - Process page table
 - Process metadata
- OS ideology
 - Multiple processes are active “simultaneously”
 - The processes have **control-flow independence** and **private address space**
- **Control-flow independence**
 - Each process executes as if the only process running on the machine
 - If there are in fact multiple processes running on the machine
 - Each process uses the CPU until pre-empted or exits
 - Then another process uses the CPU until it is pre-empted
 - Eventually the first process will get another run
 - (Effectively a priorityQ)
 - What causes pre-emption?
 - If it runs “long enough” and the OS replaces it
 - It attempts to perform a long-duration task
 - On pre-emption
 - The process’s entire **context** is saved
 - **Static information** (code, data)
 - **Dynamic state** (heap, stack, registers)
 - **OS-supplied state** (environment variables, stdin, stdout)
 - The process is flagged and suspended
 - It is placed on a process (priority) queue
 - On resuming it is restored
 - Contexts are entirely switched (**context switch**)
 - Non-static process context is held in a **process control block (PCB)**
 - Identifier (unique process ID)
 - Status (running, ready, suspended, exited)
 - State (registers)
 - Privileges (owner, group)
 - Memory management info (page table reference)
 - Accounting (CPU time used, amount done)
 - I/O (open file descriptors)
 - The OS scheduler maintains a queue
 - The OS maintains a table of PCBs
- **Unix/Linux Process**
 - Commands
 - sh - For creating processes via object-file name
 - ps - Show process info
 - w - Show per-user process info
 - top - Show high-CPU-usage process info

- kill - Send a signal to a process
- PCB info
 - pid – process id
 - ruid, euid – real and effective user id
 - rgid, egid – real and effective group id
 - Cwd
 - Accumulated execution time
 - User file descriptor table
 - Info on how to react to signals
 - Pointer to process page table
 - Process state
- Process ID (PID)
 - A positive integer, unique
 - Type pid_t
 - Process 0 is the idle process (always runnable)
 - A 'kernel artefact'
 - Not a real process
 - Needed to ensure there is always at least one process
 - Process 1 is init ("the system")
 - Low-numbered processes are system-related
 - Regular processes are in (300, maxPID (maybe 2^{16}))
- 'Families'
 - Each process had a parent process (usually that which created it)
 - A process may have child processes
 - Process 1 created at system startup
 - *If a process' parent dies, it is inherited by process 1*
- Process groups
 - Each group is associated with a unique PGID
 - Type pid_t
 - A child belongs to the process group of its parent
 - A process can create its own process group or move to another one
 - Allow the OS to keep track of groups working together
 - Signals sent to related processes
 - Management of processes for job control
 - Management of processes within pipelines
- System calls (and failure)
 - Use wrapper functions
 - Same arguments/returns as system call
 - Catches and reports the error
 - Only ever returns with a valid result
 - Not always appropriate (open() is best handled by caller)
- Unix/Linux process-related system calls
 - pid_t fork(void)
 - Requires #include <unistd.h>
 - Creates new process **by duplicating the calling process**
 - New process is the child
 - Child has different process ID


- Returns
 - In the child, fork() returns 0
 - In the parent, fork() returns the pid of the child
 - If the call fails, fork() returns -1
- Child inherits copies of parent's address space and open file descriptors
- Wrapper function for fork():

```
pid_t Fork() {
    pid_t pid;
    if ((pid = fork()) < 0) {
        perror("fork() failed");
        exit(1);
    }
    return pid;
}
```

- **void _exit(int status)**
 - Terminates current process
 - Closes any file descriptors
 - A SIGCHLD signal is sent to parent
 - Returns **status** to parent (via wait())
- **void exit(int status)**
 - Terminates current process
 - Triggers any functions registered as atexit()
 - Flushes stdio buffers, closes open FILE *s
 - Then behaves like _exit()
- **void atexit(void (*func)(void))**
 - Executes the function at exit
- **pid_t getpid() / getppid()**
 - Requires #include <sys/types.h>
 - Returns the process ID / parent process ID of the current process
 - Also getpgid(), setpgid()
- **pid_t waitpid(pid_t pid, int *status, int options)**
 - Pause current process until process pid changes state (finishing, stopping, restarting)
 - Ensures that child resources are released on exit
- **pid_t wait(int *status)**
 - Pauses until one of the child processes terminates
- **int kill(pid_t ProclD, int SigID)**
 - Requires #include <signal.h>
 - Send signal SigID to process ProclD
 - Various signals, eg:
 - **SIGHUP** – hangup detected on controlling terminal/process
 - **SIGINT** – interrupt from keyboard (ctrl+C)
 - **SIGKILL** – kill signal
 - **SIGILL** – illegal instruction
 - **SIGFPE** – floating point exception
 - **SIGSEGV** – invalid memory reference

- SIGPIPE – broken pipe
- **int execve(char *Path, char *Argv[], char *Envp[])**
 - Transforms current process by executing Path object
 - Path must be an executable, binary or script (starting with #!)
 - Passes array of strings to new process
 - Both arrays terminated by a NLL pointer element
 - Env[] contains strings of the form key=value
 - Most of the original process is lost
 - New virtual address space, signal handlers reset, etc
 - New process inherits open file descriptors from original process
- **Zombie process**
 - A process which has exited but signal not handled
 - All processes become zombie until a SIGCHLD handled
 - Parent may be delayed, but usually resolved quickly
 - Long-term zombies created when a bug in parent causes it to ignore SIGCHLD
 - Zombies occupy a slot in the process table
- **Orphan process**
 - When parent exits, orphan is given pid=1 as its parent
 - pid=1 always handles SIGCHLD when process exits
- **Process Control Flow**
 - Flow
 - Fetch instruction from memory[PC], PC++
 - Decode and execute instruction
 - If jump-type instruction, PC = new address
 - If regular instruction, carry out
 - Repeat above
 - System calls
 - Typically operate via exceptions
 - Generates an exception
 - Control transferred to exception handler
 - Carries out system-level operations
 - Returns control to the process
 - Effect is like a normal function call
 - Signals
 - Can be generated from a variety of sources
 - From another process via kill()
 - From the OS
 - From within the process
 - From a fault in the process
 - From a device
 - Processes can define how they want to handle signals
- **Signal handlers**
 - Functions invoked in response to a signal


- Knows which signal it was invoked by
- Ensures the invoking signal is blocked and carried out any appropriate action
- **SigHnd signal(int SigID, SigHnd Handler)**
 - *Defines with what function a process will handle a signal*
 - Requires <signal.h>
 - SigID is an OS-defined signal
 - Handler is one of:
 - SIG_IGN (ignore)
 - SIG_DFL (default handler)
 - A user defined function to handle SigID signals
- **int sigaction(int sigID, struct sigaction *newAct, struct sigaction *oldAct)**
 - sigID is an OS-defined signal
 - newAct defines how signal should be handled
 - oldAct saves a copy of how signal was handled
 - if newAct.sa_handler ==
 - SIG_IGN, ignore
 - SIG_DFL, default
 - struct sigaction has:
 - A pointer to a handler function
 - A mask, specifying signal to block
 - Flags to modify how signal is treated
- **Interrupts**
 - *Signals which cause normal process execution to be suspended*
 - An **interrupt handler** then carries out tasks related to interrupt
 - Control is then returned to original process
 - Egs
 - When data is fetched from the disk, interrupt and place data in a buffer
 - If a process is pre-empted
- **Exceptions**
 - *Like interrupts but for internal factors (computation)*
 - Egs
 - Division by 0
 - Failed malloc call
- **Multi-tasking summary**
 - Multiple processes “active”
 - Processes not *executing* simultaneously (although could with multiple CPUs)
 - Processes are a mixture of
 - Blocked waiting on signal
 - Runnable
 - One is running (each **CPU**)
 - Appearance of multiple simultaneous processes
 - Swap process after one runs for a defined time
 - Processes are **pre-empted**
 - **System scheduler** selects new process
- **Scheduling**
 - Processes organised into priorityQs
 - Determined by several factors:



- System processes > user processes
- Long running <
- Memory intensive <
- Some processes suggest their own priority

Device Management – *The Heavy Hardware*

- I/O devices
 - Allow programs to communicate with “the outside” world
- Memory-based data
 - Fast random access via virtual address
 - Transfer data in bytes, words, etc
- Device data
 - Must slower access
 - Random or sequential
 - Often in blocks (128B, 512B, 4KB etc)
- Hard disk
 - Address by track and sector
 - Around 0.1ms transfer
- Solid State Disk (SSD) characteristics
 - High capacity (GB), high cost, reads faster than writes
 - Pages 512B..4KB, blocks 32..128 pages, R/W page-at-a-time
 - Pages updates by erase then write
 - Average read/write time in order of 10^{-5}
- **Device Drivers**
 - Each type of device has a unique access protocol
 - Special control and data registers
 - Locations for data to be R/W
 - Device drivers – code chunks to control an i/o device
 - Core OS components
 - Sometimes written in assembler
 - Typical protocol to manipulate
 - Send request
 - Receive interrupt when completed
- Memory mapped I/O
 - OS defines special memory address
 - User programs perform i/o by getting/putting data into memory
 - Virtual addresses associated with
 - Data buffers
 - Control registers
 - Advantages:
 - Can use existing memory to access logic circuits
 - Can use full range of CPU operations
- Devices on Linux/Unix




- Device “files” can be accessed via file system under “/dev”
 - dev/diskN hard drive
 - dev/ttyN terminal device
 - dev/ptyN pseudo-terminal device
 - and also lol...
 - dev/mem the physical memory of the computer
 - dev/null data sink / empty source
 - dev/random stream of pseudo-random numbers
- **int ioctl(int FileDesc, int Request, void *Arg)**
 - Manipulates parameters of special files
 - Request is a device specific request code
 - Arg is an integer modifier or pointer to data block

Networking – *Gotta Make Those Connections*

● Sockets

- Unix/Linux provides sockets to communicate processes that may be on different systems
 - Consider them to be “a door”
 - Both endpoints need to create one
- **Socket** – an end-point of a channel
 - Either locally (Unix domain) or network-wide (internet domain)
- Server creates a socket, then
 - Binds to an address
 - Listens for connections from clients
- Client creates a socket then
 - Connects to server using an address
 - R/W to server via socket



- **int socket(int Domain, int Type, int Protocol)**
 - Requires #include <sys/socket.h>
 - Creates a socket using:
 - Domain – communications domain
 - AF_LOCAL (local host)
 - AF_INET (over the network)
 - Type – semantics of communication
 - SOCK_STREAM (sequenced, reliable communications stream)
 - SOCK_DGRAM (connectionless, unreliable packet transfer)
 - Protocol – communication protocol
 - Many exist
- **int bind(int Sockfd, SockAddr *Addr, socklen_t AddrLen)**
 - Associates open socket with an address
 - For Unix, address is a pathname
 - For Internet, address is IP and port number
- **int listen(int Sockfd, int Backlog)**

- Wait for connections on socket Sockfd
 - Allow at most Backlog connections to queue up
- **SockAddr = struct sockaddr_in**
 - sin_family – domain (AF_UNIX or AF_INET)
 - sin_port – port number
 - sin_addr – structure containing host address
 - sin_zero[8] – padding
- **accept(int Sockfd, SockAddr *Addr, socklen_t *AddrLen)**
 - Sockfd has been created, bound and is listening
 - Blocks until a connection request is received
 - Sets up a connection between client/server after connect()
- **connect(int Sockfd, SockAddr *Addr, socklen_t *AddrLen)**
 - Connects the socket Sockfd to address Addr
 - Assumes that Addr contains a process listening appropriately
- **Networks**
 - Interconnected collections of computers
 - Flavours
 - **Local area networks** (within a physical location) (**LAN**)
 - **Wide area networks** (geographically dispersed) (**WAN**)
 - **Internet** (global set of WANs)
 - Basic requirements
 - Get data from machine A to machine B
 - A and B may be separated by 100s of networks and devices
 - **How to achieve**
 - Need a unique address for destination
 - Identify a route
 - Process at intermediate nodes
 - Follow certain protocols
 - **How a file is sent**
 - File data divided into packets by source device
 - Packets are small fixed-size chunks of data with headers
 - Passed across physical link
 - Wire, radio, optic fibre
 - May not actually be “physical” physical lol
 - Passed through multiple nodes
 - Each node redirects to another node
 - Packets reach destination
 - Re-ordering, error-checking, buffering
 - File received by receiving process or user

- **The Internet**

- Concept
 - Millions of connected devices
 - Communication links
 - Fibre, copper, radio satellite
 - Packet switches
 - Routers, network switches

- **The 5-layer Model for the Internet**

- **Lowest (least abstract)**
- **Physical layer**
 - Bits
- **Link layer**
 - Ethernet, MAC addressing, CSMA
- **Network layer**
 - Routing protocols (IP)
- **Transport layer**
 - TCP – Transmission Control Protocol
 - Wait for a signal (connection established and held)
 - UDP – User Datagram Protocol
 - Keep sending, minimal checks in place, no guarantees
- **Application layer**
 - NWS, HTTP, email, FTP
- **Highest (most abstract)**

- **Protocols – effectively ‘rules and regulations’**

- Govern all communication activity on the network
 - Format and order of messages sent/received
 - Actions taken on transmission/receipt
- Defined in all the layers
 - Link: PPP (point-to-point)
 - Network: IP (internet protocol)
 - Transport: TCP (transmission control), UDP (user datagram)
 - Application: HTTP, FTP, SSH, POP, SMTP

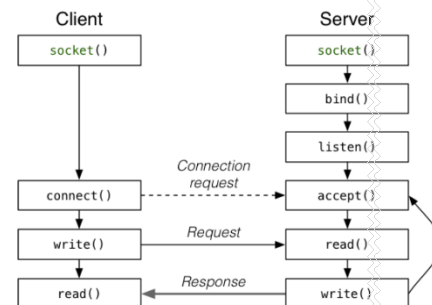

- **Client-Server Architecture**

- **Server is a data provider**
 - Process that waits for requests
 - Always-on host, permanent IP
 - Possibly using data centres / multiple CPUs
- **Client is a data consumer**
 - Sends requests to server, collects response
 - May be intermittently connected
 - Does not communicate with other clients

- **IP Addresses**

- Unique identifiers for host on a network
- Given as a 32-bit identifier (dotted quad)
 - 127.0.0.1 (refers to local host)
- IP addresses are assigned by the sys admin entering into local registry

- **The HTTP Protocol**




- HTTP = HyperText Transfer Protocol
 - Drives the web
 - Message types:
 - URLs (requests)
 - Web pages (responses)
 - Message syntax:
 - Headers and data
 - Client-server model:
 - Client = web browser
 - Server = web server

https://www.cse.unsw.edu.au:80/~cs1521/17s2/

protocol host port path

- **URL components**
 - Protocol
 - Host
 - Port
 - Path
 - Query (everything after the '?')
- Server Addresses (DNS)
 - **Domain Name System** – provides name to IP mapping
 - I.e www.cse.unsw.edu.au → IP
 - Effectively a database of mappings
 - Name servers cooperate to achieve this
 - Name resolution (two styles)
 - **Iterated query**
 - Client contacts name server X
 - Response: "I don't know, but ask name server Y" OR "Here it is"
 - **Recursive query**
 - Client contacts name server X
 - X contacts Y, Y contacts Z... until resolve
 - Management of servers
 - Top-level domain (TLD) name servers
 - Authoritative name servers (within an organisation)
 - Local (default)
 - The Unix/Linux **host** command does this too
- Transport Layer deals with
 - Data integrity
 - Some apps (file transfer) require 100% reliable transfer
 - Other apps can tolerate some loss
 - Timing
 - Some apps (eg networked games) require low transmission delay
 - Throughput
 - Some apps require minimum throughput




- Others (“elastic apps”) can use whatever
- Security
 - Some require encrypted transmission

Multiple Processes – *Threading the Needle*

- **Parallelism**
 - Multiple computations executed simultaneously
 - Eg multiple CPUs, one process on each CPU
 - Eg data vector, one processor computes each element
 - Eg SIMD (single instruction, multiple data) like above
- **Concurrency**
 - Multiple processes running (pseudo-) simultaneously
 - Eg single CPU alternating between processes
 - Can be achieved from processes
 - This requires kernel intervention
 - Each process has a significant amount of state
 - Requires system-level mechanisms
 - Increases system throughput
 - If one process is delayed, others can still run
 - Use all CPUs if multiple
 - **Effects of poorly controlled concurrency**
 - **Nondeterminism** – same code, different runs, different results
 - **Deadlock** – a group of processes wait for each other
 - **Starvation** – one process keeps missing access to a resource
 - **Other: eg bank withdrawal conundrum**
 - **Concurrency control**
 - Provide correct sequencing of interactions between processes
 - Coordinate semantically-valid access to shared resources
 - **Shared memory** based scheme
 - Uses shared variable, manipulated atomically
 - Blocks if access unavailable, decrements once available
 - **Message passing** based scheme
 - Processes communicate by sending/receiving messages
 - Receiver can block waiting for message to arrive
 - Sender may block waiting for message to be received
- **Producer-consumer problem**
 - Producer produces (puts item into an array)
 - Consumer consumes (takes item from an array)
 - We need a mechanism for a process to pause itself, and a mechanism for signalling for a process to wake up
- **Semaphores**
 - Shared memory objects to prevent these issues (literally just a shared variable)
 - Can be thought of as a count of available resources
 - **Init(Sem, InitValue), while(Sem), signal(Sem)**
 - Init – set value

- Wait – if sem>0, decrement, else wait until sem>0
 - Signal – increment sem
 - **Note: a series of Linux commands for this exist**
- **Threads** – alternate method for concurrency
 - Threads exist *within* a parent process (processes independent)
 - Threads share parent process state (processes own)
 - All threads share an address space (processes own)
 - Threads communicate via shared memory (processes use IPC)
 - Context switching between threads is cheaper than with processes
- Linux/Unix thread commands (#include <pthread.h>)
 - **int pthread_create(pthread_t *Thread, pthread_attr_t *Attr, void *(*Func)(void *), void *Arg)**
 - Creates a new thread with specified attributes
 - Thread info in *Thread
 - Starts by executing Func() with Arg
 - **pthread_t pthread_self(void)**
 - Returns pthread_t for current thread
 - Analogous to getpid()
 - **int pthread_equal(pthread_t t1, pthread_t t2)**
 - Compares two threads
 - Returns non-zero if same thread
 - **int pthread_join(pthread_t T, void **value_ptr)**
 - Suspend execution until thread T terminates
 - Exit value placed in value_ptr
 - If T has exited already, does not wait
 - **void pthread_exit(void *value_ptr)**
 - Terminate a thread and clean up
 - Store return value in value_ptr
- MapReduce
 - A model for manipulating very large data on a large network of nodes
 - Map – filter data and distribute to nodes as (key, value) pairs
 - Each node receives a set of pairs with common keys
 - Reduce – nodes perform calculation on received data items
 - Outputs from the nodes combined
- Interacting processes
 - Processes can interact via:
 - Signals, signal handlers
 - Accessing the same resource
 - Pipes (stdout of process A goes into stdin of process B)
 - Message queues
 - Sockets
 - Issues
 - Eg two processes writing to the same file ‘simultaneously’
- **File Locking**
 - **int flock(int FileDesc, int Operation)**
 - Operations:
 - LOCK_SH (acquire share lock)



- LOCK_EX (acquire exclusive lock)
 - Blocks if a process tries to acquire shared lock
 - Will fail if already locked
- LOCK_UN (unlock)
- LOCK_NB (operation fails rather than blocking)
- Does not return until lock available
- Only works if all processes accessing file use locks

• Pipes


- A style of process interaction
- Producer process writes to a byte stream, consumer process reads from same byte stream
- Provides buffered i/o between producer and consumer
 - Producer blocks when buffer full
 - Consumer blocks when buffer empty
- Pipes are bidirectional (unless processes close one file descriptor)



○ int pipe(int fd[2])

- Open two file descriptors (shared by processes)
 - fd[0] is opened for reading, fd[1] for writing
- Creating a pipe would be followed by fork() to create a child process
 - Both have copies of fd[]
 - One can write to fd[1], the other can read from fd[0]

Creating a pipe ...




○ FILE *popen(char *Cmd, char *Mode)

- Analogous to fopen()
- Cmd is passed to shell for interpretation
- Returns FILE *, which can be read/written based on Mode (or NULL if can't establish)
- Only works because both processes share an address space (both can access p[])


What if we execve() in the child? How can we still access the pipe?

- Connect stdout and p[1] in the parent process
- Connect stdin and p[0] in the child
- Then exec (parent can send data to child)
- Swap 'em for child sending data to parent
- Connecting via dup2()
 - dup2(fd1, fd2)
 - Copies file descriptor fd1 onto fd2
 - fd's are just indexes into a table of structures
 - dup(fd1) copies fd1 onto the "next free" fd



- **Message queues**

- Provides a mechanism for unrelated processes to pass information along a buffered channel shared by many processes
- Processes connect to message queues by name
- Requires `#include <mqueue.h>`
- **`mqd_t mq_open(char *Name, int Flags)`**
 - Creates a new message queue, or opens existing one
 - Flags same as `fopen()` (eg `O_RDONLY`)
- **`int mq_close(mqd_t *MQ)`**
 - Finish accessing message queue MQ
 - It continues to exist
- **`int mq_send(mqd_t MQ, char *Msg, int Size, uint Prio)`**
 - Adds message Msg to the queue MQ
 - Prio gives priority
 - If MQ is full, blocks until MQ space available
- **`mq_timedsend()`**
 - Waits for specified time if MQ full
 - Fails if still no space on MQ after timeout



CRAFT

Bits and Bytes – *Delicious!*


- One byte is 8 bits
 - `int` = 4 bytes = 32 bits
- C bitwise operators
 - **`&` (AND)**
 - Goes through every digit and applies the operator to each matching digit
 - **`|` (OR)**
 - Likewise
 - **`^` (XOR)**
 - Result is 1 if the compared numbers are different, else 0
 - **`~` (NEG)**
 - Negation – converts all zeroes to ones and ones to zeroes
 - **`<<` (left shift)**
 - Sends digits to the left by a given amount and replaces with 0s
 - Eg: “01011011 << 3” gives 11011000
 - **`>>` (right shift)**
 - Sends digits to the right by a given amount and replaces with 0s**
 - **unless the quantity is signed, then it is replaces with the first digit
 - Eg: “01011011 >> 3” gives 00010110

- Representing numbers
 - Hexadecimal**
 - Each number is 4 bits
 - Eg 4 = 0100
 - Eg F = 1111
 - When writing hex in C, begin it with 0x
 - Eg 0x71 for 113
 - Octal**
 - Each number is 3 bits
 - Eg 5 = 101
 - Eg 6 = 110
 - Binary**
 - Each digit is one bit
 - Eg 1 = 1
 - Eg 0 = 0
 - When writing binary in C, begin it with 0b
 - Eg 0b1101 for 13

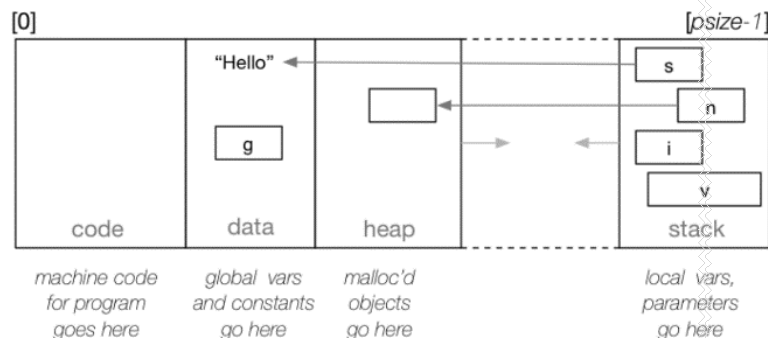
Data Representation – *Let Me Lay it All Out for Ya*

- Typical memory usage for a given C code
 - Space in between malloc and stack
 - Stack needs the space, keeps adding variables 'on top' as required by the code
 - Different "stack frames" for different 'levels' of the stack
 - Eg if we call a function in a function, we put another frame on the stack
 - Malloc is fluid and so needs its distance

```
int main() {
    int n, m;
    n = 5; m = f(n);
}
int f(int x) {
    return g(x);
}
int g(int y) {
    int r = 4 * h(y);
    return r;
}
int h(int z) {
    int i, p = 1;
    for (i=1; i<=z; i++)
        p = p * i;
    return p;
}
```



- The C view of data
 - Variables are an example of computational objects
- Computational objects** have:
 - A **location** in memory (obtain with '&')
 - A **value** (a bit string)
 - A **name** (unless malloc'd)
 - A **type**, which determines
 - Size** (in units of whole bytes, found with sizeof)
 - Which **options** to use to **interpret** its value
 - A **scope** (where it's visible in the program)
 - A **lifetime** (during which part of execution it exists)
 - Egs
 - A global variable has global region, scope of the whole file, lifetime of all runtime




- A local variable is in the stack, has a scope in main(), and has a lifetime ending with main()
- Static variables in main have global region, scope in main() and lifetime of the whole runtime
- Variables defined *after* main() can *only* be accessed in functions
- Functions are in the code, with scope everywhere and lifetime of their execution
- Function arguments are in the stack, with scope and lifetime in f()
 - Same for local function variables

- The physical view of data – essentially a very large array of bytes

- Called main memory, RAM, primary storage, etc
 - Indexes are memory addresses (pointers)
 - Data can be fetched in chunks of 1,2,4,8 bytes
 - The cost of fetching any byte (no matter where it is) is the same (nanoseconds)
 - The byte address for an N-byte object must be divisible by N
 - Can be volatile or non-volatile

- **-endian**

- **Little-endian**
 - Least significant byte first, then second least second, etc
- **Big-endian**
 - “It’s when you look at it and it looks nice.” -Annie
 - Most significant byte first, then second second, etc



- Data representation

- **ASCII (ISO 646)**

- 7-bit values, using lower 7-bits of a byte, with the top bit 0

- **UTF-8 (Unicode)**

- 8-bit values, with ability to extend further
- Can encode all human language and other symbols, including mathematical signs or emojis
- Uses a variable length encoding (one, two, three or four bytes)
- The first byte begins with a code instructing the compiler to detect ASCII

#bytes	#bits	Byte 1	Byte 2	Byte 3	Byte 4
1	7	0xxxxxxx	-	-	-
2	11	110xxxxx	10xxxxxx	-	-
3	16	1110xxxx	10xxxxxx	10xxxxxx	-
4	21	11110xxx	10xxxxxx	10xxxxxx	10xxxxxx

The 127 1-byte codes are compatible with ASCII
 The 2048 2-byte codes include most Latin-script alphabets
 The 65536 3-byte codes include most Asian languages
 The 2097152 4-byte codes include symbols and emojis and ...

- Eg μ is U+00B5 which is 1011 0101 in binary and unicode **11000010 10110101**

- Integers (three ways to write)

- Signed decimal 42 (0-9)
- Unsigned hexadecimal **0x2A** (0-F)
- Signed octal **052** (0-7)
- Variations:
 - Unsigned int 123U (32-bit)
 - Long int 123L (64-bit)
 - Short int 123S (16-bits)
- Easy data conversions
 - Binary \rightarrow hex: look in blocks of four digits

- **Signed integers**

- The first bit (digit) is 0 for positive, 1 for negative
 - Several representation options
- **Signed magnitude**
 - First digit is sign, the rest are magnitude
 - Simple addition does not work
- **Ones complement**
 - -N is formed by converting all bits in N
 - But we have two zeros: 11111111 (-0) and 00000000 (0)
- **Twos complement**
 - -N is formed by converting all bits in N and adding 1
 - $x = -(-x)$
 - Addition works
 - There's a random 'overflow' bit we conveniently ignore

- **Different types of pointers to point to different things**


- However they all have the same size (typically 32-bit)
- The values must all be appropriate for the given data point
 - (char *) can reference any byte address
 - (int *) must have $\text{addr}\%4 == 0$

- **Considering the diagram**

- Low memory values tend to be associated with the code
- High memory values tend to be associated with the stack

- **Incrementing pointers (pointer arithmetic)**

- Moves them to the next memory address forward
- A common (efficient) method for scanning a string is to increment a character pointer



- **Normalisation** – “New binary scientific notation”

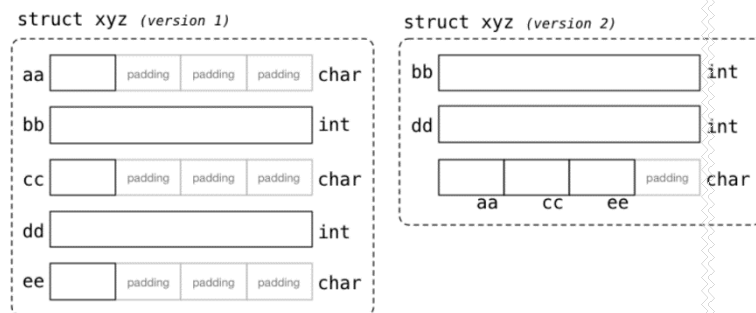
- Floats are represented by: $\text{Num} = (1 + \text{fraction}) * 2^{(\text{exp} - \text{bias})}$
- Exponent is represented relative to a bias value B
 - $\text{Bias} = 2^{n\text{bits}-1} - 1$
- 1010.1011 is normalised as 1.0101011×2^{011}
- Always have 1 before the decimal point
 - If we had a zero we'd be wasting memory, we could just shift the whole thing

- **Floating point**

- Float – typically 32 bit
 - 1 bit for sign
 - 8 bits for exponent
 - 23 bits for fraction
 - Each bit is $+2^{-1}$, $+2^{-2}$, etc
- Double – typically 64 bit
 - 1 bit for sign
 - 11 bits for exponent
 - 52 bits for fraction

- Example
 - 0 10000000 110000000000000000000000
 - Sign is positive
 - Exponent is $128 - Bias = 1$
 - We thus multiply the fraction by $2^{exp} = 2^1$
 - Fraction is $\frac{1}{2} + \frac{1}{4} = 0.75$
 - Number is thus 1.75×2^1
- Representation
 - `printf("%X.Ylf", number)`
 - Prints X digits total (including the DP)
 - Y of them are after the DP

- Structs
 - Order of struct elements is largely determined by compiler
 - Re-ordering fields may save space



The Command Line – *A Strange Place*

- Compiling
 - `gcc -E x.c` (executes the C pre-processor, rewriting code)
 - `gcc -S x.c` (produces a file x.s containing assembly code)
 - `gcc -c x.c` (produces a file x.o containing relocatable machine code)
 - `gcc x.c` (produces a.out)
 - `gcc *.o` (compiles everything)
- Makefiles
 - Sometimes you don't want to compile every file, but keeping track of which files you want to compile is tedious
 - Between files:
 - Dependencies exist
 - Eg x.o depends on x.c and y.h
 - Actions state how to produce targets from sources
 - Rules combine dependencies and actions
 - Eg if x.c changes, rebuild x.o
 - Makefiles contain definitions and rules for compilation
 - Target
 - Source
 - If you change one of these, the blue and red files will recompile upon 'make'
 - Action
 - Other arguments
 - CC = compiler

```
bm : bm.o Stack.o
    gcc -o bm bm.o Stack.o

bm.o : bm.c Stack.h
    gcc -c -Wall -Werror bm.c

Stack.o : Stack.c Stack.h
    gcc -c -Wall -Werror Stack.c
```

- CFLAGS = (compiler flags)
 - -Wall (display all warnings)
 - -Werror (display all errors)
- Commands
 - cat file
 - Prints the contents of the file
 - man some_command
 - Opens the manual page for that command

MIPS – *Man, I Program Sowell*

- A well-known and relatively simple architecture
 - Silicon graphics, NEC, Nintendo64, etc
- MIPS features
 - 32 x 32-bit general purpose registers
 - 16 x 64-bit double precision registers
 - PC – 32-bit register
 - HI, LO – store the results of multiplication and division
 - Two 32-bit registers to store big integers
- Data types
 - Byte (8 bits)
 - Halfword (2 bytes)
 - Word (4 bytes)
 - Characters require 1 byte of storage
 - Integers require 1 word of storage
- Reference
 - Registers can be referred to as \$0,...,\$31 or by symbolic names

Register	Name	Notes	Cont
\$0	zero	The value 0; unchangeable	
\$1	\$at	Assembler temporary; reserved for assembler use	
\$2	\$v0	Value from expression evaluation or function return	This is what we use for all syscalls. If we load an integer it goes here. Also used for function returns.
\$3	\$v1	Value from expression evaluation or function return	
\$4	\$a0	First argument to a function/subroutine, if needed	This mean when we call a function, it immediately looks to the \$a? registers
\$5	\$a1	Second ...	
\$6	\$a2	Third ...	
\$7	\$a3	Fourth ...	

\$8,...,\$15	\$t0,...,\$t7	Temporary; must be saved by caller to subroutine; subroutine can overwrite	
\$16,...,\$23	\$s0,...,\$s7	Safe function variable; must not be overwritten by called subroutine	We modify these within functions
\$24, \$25	\$t8, \$t9	Temporary; must be saved by caller to subroutine; subroutine can overwrite	
\$26, \$27	\$k0, \$k1	For kernel use; may change unexpectedly	
\$28	\$gp	Global pointer	
\$29	\$sp	Stack pointer	
\$30	\$fp	Frame pointer	
\$31	\$ra	Return address of most recent caller	We store the pointer to the place in our program whence we jumped.

- There are also floating point registers
- Features of MIPS language
 - # comments
 - Labels appended with ‘:’
 - Directives – a symbol beginning with ‘.’
 - **Directives list:**
http://students.cs.tamu.edu/tanzir/csce350/reference/assembler_dir.html
 - Assembly language instructions
- **Pre-main tags:**
 - .data
 - .data on a line on its own introduces the data section
 - Variables can now be introduced
 - a: .word 42
 - array: .space 20
 - msg: .asciiz “Hello world”
 - Etc
 - .globl
 - .globl main declares the main function
 - .text
 - .text on a line on its own introduces the text section
 - Functions can now be introduced (beginning with their own tag)
 - main:

- General code structure and example code

```
# Prog.s ... comment giving description of function
# Author ...

.data      # variable declarations follow this line
           # ...

.text      # instructions follow this line
.globl main
main:      # indicates start of code
           # (i.e. first user instruction to execute)
           # ...

# End of program; leave a blank line to make SPIM happy
```

```
.data
a: .word 42      # int a = 42;
b: .space 4      # int b;
.text
.globl main
main:
    lw  $t0, a      # reg[t0] = a
    li  $t1, 8      # reg[t1] = 8
    add $t0, $t0, $t1 # reg[t0] = reg[t0]+reg[t1]
    li  $t2, 666    # reg[t2] = 666
    mult $t0, $t2    # (Lo,Hi) = reg[t0]*reg[t2]
    mflo $t0        # reg[t0] = Lo
    sw  $t0, b      # b = reg[t0]
    ....
```

- Instructions (all 32-bit)
 - **Instructions list:**
http://www.dsi.unive.it/~gasparetto/materials/MIPS_Instruction_Set.pdf
 - lw – load word (assign value from a memory (word) address)
 - li – load integer (assign value)
 - la – load address
 - add – add (requires three variables)
 - mult – multiply (requires two variables and stored in Lo,Hi)
 - mflo / mghi – store the value from Lo / Hi in a specified variable (must call one at a time)
 - sw – store word (store a value into a memory (word) address)
 - mov – move a value from one register to another
- Memory layout

Region	Address	Notes
Text	0x00400000	Contains only instructions (read only)
Data	0x10000000	Data objects (read/write)
Stack	0x7fffefff	Grows down (read/write)
k_text	0x80000000	Kernel code. Only accessible kernel mode (read only)
k_data	0x90000000	Kernel data. Only accessible kernel mode (read only)



- **Classes of instructions**
 - Load and store
 - Transfer data between registers and memory
 - Computational
 - Performs arithmetic/logical operations
 - Jump and branch
 - Transfer control of program execution
 - Coprocessor
 - Standard interface to various co-processors
 - Special
 - Miscellaneous tasks

- **syscall**

- A special instruction
- Several types of calls
- Called from a \$v register
 - li \$v0, 5
syscall
 - The above scans a value into \$v0

Service	Code	Arguments	Result
print_int	1	\$a0 = integer	
print_float	2	\$f12 = float	
print_double	3	\$f12 = double	
print_string	4	\$a0 = char *	
read_int	5		integer in \$v0
read_float	6		float in \$f0
read_double	7		double in \$f0
read_string	8	\$a0 = buffer, \$a1 = length	string in buffer (including "\n0")

- SPIM instructions:



<https://cgi.cse.unsw.edu.au/~cs1521/19T2/docs/spim.php>

- SPIM has its own useful instructions that are mapped to one of more MIPS instructions
- Directive note
 - Directives apply to everything below the directive, before a new directive
- The importance of .align
 - .align n means we jump to the nearest specified block of memory dividible by 2^n
 - This is critical for storing certain data types
- Addressing modes
 - lw \$t0, var (address via naming)
 - lw \$t0, (\$s0) (indirect addressing)
 - lw \$t0, 4(\$s0) (indexed addressing – accesses from the fourth element/byte)

```
lw    $t1,label    # reg[t1] = memory[&label]
sw    $t3,label    # memory[&label] = reg[t3]
                    # &label must be 4-byte aligned
la    $t1,label    # reg[t1] = &label
lui   $t2,const    # reg[t2] = const << 16
and   $t0,$t1,$t2  # reg[t0] = reg[t1] & reg[t2]
add   $t0,$t1,$t2  # reg[t0] = reg[t1] + reg[t2]
                    # add as signed (2's complement) ints
addi  $t2,$t3, 5   # reg[t2] = reg[t3] + 5
                    # "add immediate" (no sub immediate)
mult  $t3,$t4      # (Hi,Lo) = reg[t3] * reg[t4]
                    # store 64-bit result in registers Hi,Lo
seq   $t7,$t1,$t2  # reg[t7] = (reg[t1]==reg[t2]) ? 1 : 0
j     label        # PC = &label
beq   $t1,$t2,label # PC = &label if (reg[t1] == reg[t2])
nop
```

```
.text    # following instructions placed in text
.data    # following objects placed in data

.globl   # make symbol available globally

a: .space 18    # uchar a[18]; or uint a[4];
   .align 2     # align next object on 2^2-byte addr

i: .word 2      # unsigned int i = 2;
v: .word 1,3,5  # unsigned int v[3] = {1,3,5};
h: .half 2,4,6  # unsigned short h[3] = {2,4,6};
b: .byte 1,2,3  # unsigned char b[3] = {1,2,3};
f: .float 3.14  # float f = 3.14;

s: .asciiz "abc"
   # char s[4] {'a','b','c','\0'};
t: .ascii "abc"
   # char s[3] {'a','b','c'};
```

- Like pointer arithmetic, but this only does 4 bytes (*not* 4*4 bytes)
- **Addresses** are hard :((a and b are registers and/or variables)
 - `la a, b`
 - Takes the ADDRESS of b and puts that in a
 - `lw a, b`
 - Takes the CONTENT of b and puts that in a
 - `li a, num`
 - Puts the number num into a
 - `lw a, (b)`
 - $a = *b$
 - `lw a, 4(b)`
 - $a = *b + 4$
- **Additional loading**
 - `lb` load one byte (`lbu` for unsigned)
 - `lh` load two bytes (`lhu` for unsigned)
 - `lw` load four words
 - `la` load the address

- **While loops**

```
while ($s1 < $s0) {
    ...
}

while:
    bge $s1, $s0, end_while
    ...
    addi $s1, $s1, 1

end_while:
```

More C to MIPS:

<code>int x = 5;</code>	<code>x: .word 5</code>
<code>int y;</code>	<code>y: .space 4</code>
<code>z = 5 * (x+y);</code>	<code>t = x + y;</code> <code>lw \$t0, x</code> <code>lw \$t1, y</code> <code>add \$t0, \$t0, \$t1</code>
	<code>t = 5 * t;</code> <code>li \$t1, 5</code> <code>mul \$t0, \$t0, \$t1</code>
	<code>z = t;</code> <code>sw \$t0, z</code>
	<code>x = 2;</code> <code>li \$t0, 2</code> <code>sw \$t0, x</code>

```

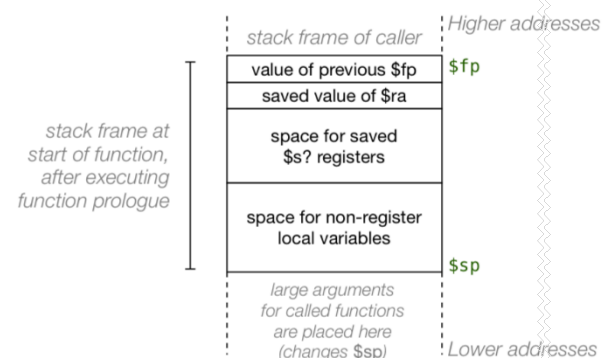
i = 0; n = 0;      i = ; n = 0;
while (i < 5) {    loop:
    n = n + i;      if (i >= 5) goto end;
    i++;            n = n + i;
}                  i++;
                   goto loop:
                   end:

if (cond)          if_stat:          if_stat:
    {statements1}   t0 = (cond)        t0 = evaluate (cond)
[not actual MIPS]
else               if (t0 == 0)        beqz $t0, else_part
    {statements}    goto else_part     statements1
                   statements1        j end_if
                   goto end_if        else_part:
                   else_part:         statements2
                   statements2        end_if:
                   end_if:

```

• Functions

- The process
 - The arguments are evaluated and set up for function
 - Control is transferred to the code for the function
 - Local variables are created
 - Function code executes
 - The return value is set up
 - Control transfers back to where the function was called from
 - The caller receives the return value
- Stack frames
 - Each function allocates a stack frame
 - Uses:
 - Saved registers
 - Local variables
 - Parameters to callees
 - Created in the *prologue* (push)
 - Removed in the *epilogue* (pop)
- Conventions
 - Caller saved registers (\$t0...\$t9, \$a0...\$a3, \$ra)
 - f() calls g() knowing these registers may change
 - Callee saved registers (\$s0...\$s7, \$sp, \$fp)
 - f() calls g() assuming the registers will be unchanged
 - If g() modifies them, they will first be saved and then restored
- Call with instruction "jal func"
- Put return value in \$v0



- End instruction is “jr \$ra”
 - This jumps back to where we left off with jal

```
# Set up stack frame.
sw  $fp, -4($sp)    # push $fp onto stack
la  $fp, -4($sp)    # set up $fp for this function
                        # Using $sp and $fp we chuck all our
                        # variables onto the stack
sw  $ra, -4($fp)    # save return address
sw  $s0, -8($fp)    # save $s0 to use as ...
sw  $s1, -12($fp)   # save $s1 to use as ...
sw  $s2, -16($fp)   # save $s2 to use as ...
addi $sp, $sp, -20  # move $sp to top of stack

# Alternatively:

addi $sp, $sp, -4
sw  $fp, ($sp)
move $fp, $sp
addi $sp, $sp, -4
sw  $ra, ($sp)
addi $sp, $sp, -4
sw  $s0, ($sp)
addi $sp, $sp, -4
sw  $s1, ($sp)
addi $sp, $sp, -4
sw  $s2, ($sp)

# clean up stack frame
                        # We retrieve all our variables
                        # from the stack
lw  $s2, -16($fp)   # restore $s2 value
lw  $s1, -12($fp)   # restore $s1 value
lw  $s0, -8($fp)    # restore $s0 value
lw  $ra, -4($fp)    # restore $ra for return
la  $sp, 4($fp)     # restore $sp (remove stack frame)
lw  $fp, ($fp)      # restore $fp (remove stack frame)

jr  $ra
```

- **Function calling protocol**

- Before calling
 - Place arguments in \$a0...\$a3
 - If more than 4 args, push all args onto the stack

- Save any non-\$s registers that need to be preserved by pushing onto the stack
- jal address of the function (usually labelled)
- Pushing onto the stack (eg \$t0)
 - addi \$sp, \$sp, -4
sw \$t0, (\$sp)
 - OR
 - sw \$t0, -4(\$sp)
addi \$sp, \$sp, -4
-

```
# start of function
FuncName:
# function prologue
# set up stack frame ($fp, $sp)
# save relevant registers (incl. $ra)
...
# function body
# perform computation using $a0, etc.
# leaving result in $v0
...
# function epilogue
# restore saved registers (esp. $ra)
# clean up stack frame ($fp, $sp)
jr $ra
```

• Arrays

- a: .space 20 (holds 5 elements, ie 20 bytes worth of words)
- a: .word 1, 3, 5, 7, 9 (a = {1,3,5,7,9})
- **How do we access the numbers?**
- *Increment the memory location by 4 (if we are dealing with words)!*

```
int sum, i;          sum: .word 4      # use reg for i
int a[5] = {1,3,5,7,9}; a:  .word 1,3,5,7,9
...
sum = 0;              ...
for (i = 0; i < N; i++) for: bgt $t0, $t2, end_for
                        move $t3, $t0
                        mul  $t3, $t3, 4
                        add  $t1, $t1, a($t3)
                        addi $t0, $t0, 1 # i++
                        j     for
                        end_for: sw  $t1, sum
                                move $a0, $t1
                                li  $v0, 1
                                syscall # printf
```

- Note: a(\$t3) is &a + \$t3

• 2D Arrays

- To access array[row][col]
- We access offset(array), where offset = row*rowsize + col
- **Make sure to multiply this offset by 4 if we are dealing with words!**

• Structs

- Literally just define space
- Stu1: .space 56
- You then put values into the struct according to offsets

- You save pointers to structs if you want to work with them
- **argc and argv**
 - These are just put on main's stack frame
- **Data structures**
 - char as one byte in memory
 - int as one word in memory
 - double as two words in memory
 - All of the above could be implemented:
 - In a register (small scope)
 - On the stack
 - In .data if we need longer lifetimes (but don't do this lol)
 - arrays sequence of memory bytes/words
 - structs chunk of memory accessed by offsets
 - linked structures struct containing address of another struct
- Static VS dynamic
 - Static:
 - Uninitialized memory is allocated at compile time
 - val: .space 4
 - Initialised memory allocated at compile time
 - val: .word 5
 - Dynamic:
 - Put things in registers